

# More advanced C

Memory allocation, Structs  
Files, Directories, and I/O



# Memory Allocation

# Dynamic memory allocation in a nutshell

- Dynamic memory allocation = request memory and free it when done
  - New allocations are on the Heap
- It's up to the user to keep track of that
  - Pointers keep track of addresses to such requested memory chunks
- **Why do we need it?**
  - To expand and shrink data structures dynamically

# Dynamic data structures

- Can you think of any examples?




# Example: multiplayer games




# Game map representation

- Goal: fast retrieval of game objects
- Game map representation: tree structure




# Sample gamemap tree



Game map


Split the game map into regions based on number of players




Tree representation  
for fast access

Each player is stored in a list in the corresponding tree node

# Sample gamemap tree



Game map



Tree representation  
for fast access

Tree grows and shrinks  
based on player movements




# Dynamic memory allocation

## - summary -

- We don't know in advance how much memory our data structure will occupy
- Need to be able expand and shrink data structures during program execution
- Mechanism to request for a dedicated new chunk of memory
  - Malloc() => returns a pointer containing the address of the new chunk!
  - Free() => relinquishes the memory chunk!



AND SUDDENLY YOU  
MISSTEP, STUMBLE,  
AND JOLT AWAKE?



xkcd.com

# Static Allocation

- Recall: static allocation happens at compile time based on variable definitions.

```
int x = 2;  
int a[4];  
int *b;
```

```
int main() {}
```

## SYMBOL TABLE:

main	0x804837c	.text	f9
x	0x8049588	.data	04
b	0x8049688	.bss	04
a	0x804968c	.bss	10

0x804837c *main*

0x804957c *init.data*

0x8049588 2

0x8049684 *uninit. data*

0x8049688 ???

0x804968c ???

0x8049690 ???

0x8049694 ???

0x8049698 ???

# Dynamic Memory Allocation

- In Java,

```
Set s; // Memory is allocated for pointer s
// Memory is allocated for object
s = new HashSet();
```

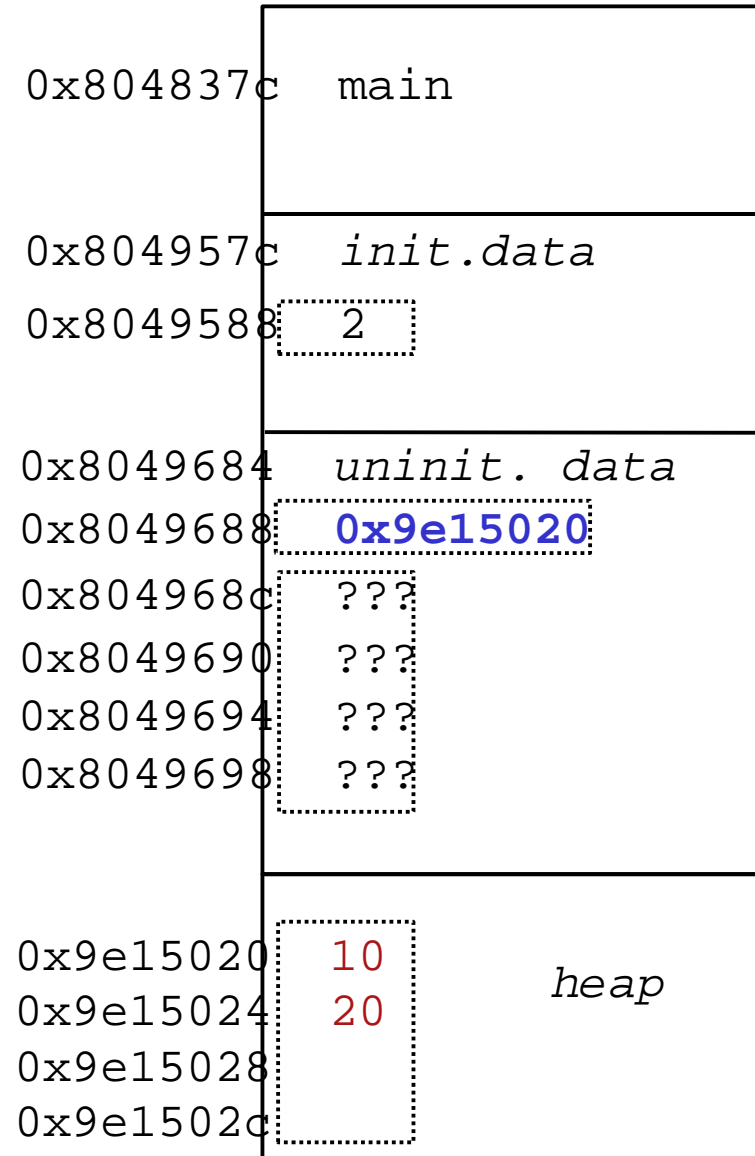
- In C,

```
int *a; /* Memory is allocated for pointer a */
/* Memory is allocated for a to point to */
a = malloc(10 * sizeof(int));
```

# Dynamic Allocation

```
int x = 2;
int a[4];
int *b;

int main() {
    b = malloc(4 *
sizeof(int));
    b[0] = 10;
    b[1] = 20;
}
```



# Always check manual pages

\$ man malloc

## SYNOPSIS

```
#include <stdlib.h>
```

```
void *calloc(size_t nmemb, size_t size);
```

```
void *malloc(size_t size);
```

```
void free(void *ptr);
```

```
void *realloc(void *ptr, size_t size);
```

## DESCRIPTION

**malloc()** allocates size bytes and returns a pointer to the allocated memory. The memory is not cleared.

**free()** frees the memory space pointed to by ptr, which must have been returned by a previous call to **malloc()**, **calloc()** or **realloc()**. Otherwise, or if **free(ptr)** has already been called before, undefined behaviour occurs. If ptr is **NULL**, no operation is performed.

# malloc

```
void *malloc(size_t size);
```

- Some things you haven't seen yet:

`void *`

- A generic pointer type that can point to memory of any type.

`size_t`

- A type defined by the standard library as the type returned by `sizeof`.
- The type is `unsigned long`.

# malloc

- Can always assign a void pointer to any more specific type of pointer.

```
int *i = (int*)malloc(sizeof(int)); //  
    type cast not mandatory
```

```
int *i = malloc(sizeof(int));      //  
    implicit conversion
```

```
char *c = malloc(NAME_SIZE);
```

- sizeof works on types, and knows type of expressions.

```
double *d = malloc(5*sizeof(*d));
```

- Be careful to allocate the correct number of bytes.
- E.g., `int *i = malloc(1); /*wrong*/`
  - allocates 1 byte, not 1 int.



# NULL pointers

- A function that returns a block of memory might fail to do so, in which case it returns a NULL pointer.
- NULL is a pre-processor variable defined in `stdlib.h` (included from `stdio.h`) and other places
  - it is usually defined to be 0 (no program allocates anything at address 0x0)

# De-allocating memory

```
int *a = malloc(10 * sizeof(int));  
int b[10];  
...  
a = b;
```

- What is wrong with the last line? It compiles and runs fine.
- We have lost the pointer to the memory region allocated in the first line, so that space is now tied up until the program terminates.

⇒ **Memory leak!**

# free()

- Before removing the last pointer to a memory region, you must explicitly deallocate it.
  - No garbage collection in C!

```
int *a = malloc(10 * sizeof(int));  
int b[10];  
...  
free(a);  
a = b;  
/*No memory leak */
```

Is "a" NULL after the free statement?

→ No, `free` cannot change the value of a parameter

# Dangling pointers

```
int *a = malloc(10 * sizeof(int));  
...  
free(a);  
printf("%d\n", a[0]); /* Error */
```

- Dereferencing a pointer after the memory it refers to has been freed is called a “dangling pointer”.
- Behaviour is undefined. Might:
  - appear to work
  - bogus data
  - program crash

# Dangling pointers

```
int *a = malloc(10 * sizeof(int));  
...  
free(a);  
printf("%d\n", a[0]); /* Error */
```

- Can you re-use pointer “a” after free() though?
  - Yes, recall that memory to store the pointer is allocated by default; However, allocating memory for the location where the pointer is pointing to, that’s up to the programmer!

```
free(a); /* memory is released => a points to an invalid  
location */  
  
int i = 5;  
a = &i; /* valid, now points to the address of variable i */  
printf("%d\n", *a);  
  
a = malloc(15 * sizeof(int)); /* valid, we request new memory */  
printf("%d %d\n", a[0], a[5]);
```

# Arrays of pointers

- Most obvious use is to get an array of strings
  - Consider a word = an array of chars
  - A sentence = an array of words

```
#define LEN 4                // define macro, constant
char **strs = malloc(3*sizeof(char *)); // define an array of 3
words

for(i = 0; i < 3; i++) {
    strs[i] = malloc(LEN);    // each word itself has to be
    allocated
}
strs[0] = strncpy(strs[0], "209", LEN); // copy a word in strs[0]
strs[1] = strncpy(strs[1], "369", LEN); // copy a word in strs[1]
```

- What else can we represent?
  - A matrix :    int \*\*a;  
  (static allocation:    int a[10][10]; )
  - An array of matrices:    int \*\*\*a;  
  (static allocation:    int a[10][10][5]; )

# Copying or moving memory

```
$ man memcpy
```

```
NAME
```

```
memcpy - copy memory area
```

```
SYNOPSIS
```

```
#include <string.h>
```

```
void *memcpy(void *dest, const void *src, size_t n);
```

```
DESCRIPTION
```

The **memcpy()** function copies `n` bytes from memory area `src` to memory area `dest`. The memory areas must not overlap. Use **memmove(3)** if the memory areas do overlap.

# Copying or moving memory

- What's the difference between these 2:

```
int *p, *q, i;  
p = malloc(10*sizeof(int));  
q = malloc(20*sizeof(int));  
for (i = 0; i < 10; i++)  
    p[i] = i;
```

1) `q = p; // what also happens here?`

OR

2) `memcpy(q, p, 10*sizeof(int));`



# Tips

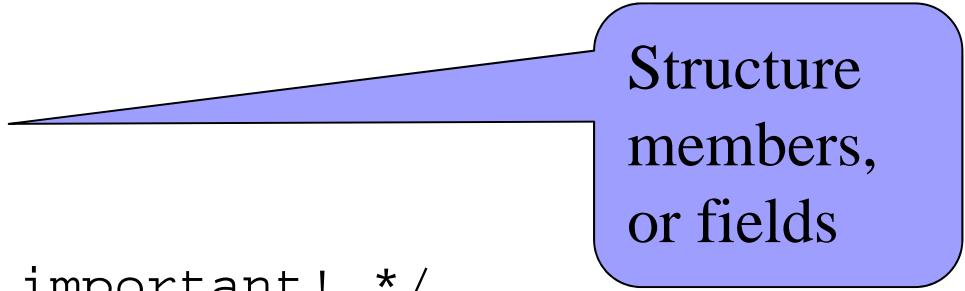
- Use a debugger and start to figure out what valid addresses look like.
- Check return values from library functions.
- Watch out for common errors:
  - forgetting to allocate memory when a pointer is declared
  - dereferencing a pointer after it's been free'd
  - losing track of a memory block without free'ing (memory leak)!
- Remember: practice, practice, practice!

# Structures

# Structs

- A collection of related data items

```
struct Point {  
    float x;  
    float y;  
};  
/* Semicolon is important! */
```



Structure  
members,  
or fields

```
struct Point p;  
p.x = 1.53;  
p.y = 8.27;
```

```
struct Point q;  
q.x = p.x;  
q.y = p.y;
```

# Structs

```
struct student {  
    char *name;  
    int age;  
};
```

```
struct student s1; /*allocates space for the record */  
s1.name = malloc(4*sizeof(char));/* don't forget this! */  
s1.name[0] = 'J';  
s1.name[1] = 'O';  
s1.name[2] = 'E';
```

- Pointers use '->' instead of '.' to refer to struct members!

```
struct student *s2;  
s2 = malloc(sizeof(struct student)); /* allocate pointer */  
s2->name = malloc(4*sizeof(char));/* again, don't forget this!*/
```

- To simplify syntax – use typedef:

```
typedef struct student Student;  
Student s3, *s4;  
s4 = malloc(sizeof(Student));
```

# Structs as arguments

```
/* Remember: pass-by-value */
void print_student(struct student s) {
    printf("Name = %s\n", s.name);
    printf("Age   = %d\n", s.age);
}

int main() {
    struct student s1, *s2;
    ...
    print_student(s1);
    print_student(*s2);
}
```

# Passing pointer or struct?

```
/* Incorrect */  
void incr_age(struct student *r) {  
    r.age++;  
}
```

```
/* Correct */  
void incr_age(struct student *r) {  
    r->age++;  
}
```

# Concrete Example

```
int stat(const char *file_name, struct stat *buf);
```

```
struct stat {  
    dev_t      st_dev;      /* device */  
    ino_t      st_ino;      /* inode */  
    mode_t     st_mode;     /* protection */  
    nlink_t    st_nlink;    /* number of hard links */  
    uid_t      st_uid;      /* user ID of owner */  
    gid_t      st_gid;      /* group ID of owner */  
    dev_t      st_rdev;     /* device type (if inode device)*/  
    off_t      st_size;     /* total size, in bytes */  
    blksize_t  st_blksize;  /* blocksize for filesystem I/O */  
    blkcnt_t   st_blocks;   /* number of blocks allocated */  
    time_t     st_atime;    /* time of last access */  
    time_t     st_mtime;    /* time of last modification */  
    time_t     st_ctime;    /* time of last change */  
};
```

# stat

- By calling the `stat` function on a filename you want to fill in the fields of the `struct stat`.
- You must pass in a pointer, and there must be space allocated!!!

```
struct stat sbuf;  
if(stat("myfile", &sbuf) == -1) {  
    perror("stat");  
    exit(1);  
}  
printf("Owner = %d", sbuf.st_uid);
```



# Common error

```
struct stat *sbuf;  
if(stat("myfile", sbuf) == -1)  
{  
    perror("stat");  
    exit(1);  
}
```

# NEXT UP

- Files, I/O
- Strings